

```
In [2]: !pip install transformers[sentencepiece] datasets sacrebleu rouge_score py7zr -q
```

```
0[K] | 5.5 MB 39.2 MB/s
0[K] | 451 kB 71.7 MB/s
0[K] | 118 kB 89.5 MB/s
0[K] | 65 kB 4.4 MB/s
0[K] | 182 kB 89.7 MB/s
0[K] | 115 kB 87.5 MB/s
0[K] | 212 kB 75.9 MB/s
0[K] | 127 kB 87.1 MB/s
0[K] | 357 kB 95.7 MB/s
0[K] | 50 kB 6.7 MB/s
0[K] | 2.3 MB 74.0 MB/s
0[K] | 379 kB 87.1 MB/s
0[K] | 138 kB 76.8 MB/s
0[K] | 93 kB 2.4 MB/s
0[K] | 7.6 MB 78.7 MB/s
0[K] | 1.3 MB 89.9 MB/s
0[?25h Building wheel for rouge-score (setup.py) ... 0[?25l0[?25hdone
```

```
In [3]: !nvidia-smi
```

```
Tue Nov 29 04:03:54 2022
+-----+
| NVIDIA-SMI 460.32.03      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+
| GPU   Name               Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap     |      Memory-Usage | GPU-Util  Compute M. |
|                                           | MIG M.         |
+-----+-----+
|  0   Tesla T4            Off          | 00000000:00:04:0 Off |                    0 |
| N/A   39C    P8             9W /  70W |  0MiB / 15109MiB |      0%      Default |
+-----+-----+

Processes:
+-----+
| GPU   GI    CI        PID   Type   Process name                      GPU Memory |
| ID    ID                                 |                 Usage                   |
+-----+-----+
| No running processes found          |
+-----+
```

```
In [4]: from transformers import pipeline, set_seed

import matplotlib.pyplot as plt
from datasets import load_dataset
import pandas as pd
from datasets import load_dataset, load_metric

from transformers import AutoModelForSeq2SeqLM, AutoTokenizer

import nltk
from nltk.tokenize import sent_tokenize

from tqdm import tqdm
import torch

nltk.download("punkt")
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
```

```
Out [4]: True
```

```
In [5]: from transformers import AutoModelForSeq2SeqLM, AutoTokenizer

device = "cuda" if torch.cuda.is_available() else "cpu"

model_ckpt = "google/pegasus-cnn_dailymail"

tokenizer = AutoTokenizer.from_pretrained(model_ckpt)

model_pegasus = AutoModelForSeq2SeqLM.from_pretrained(model_ckpt).to(device)
```

```
Downloading: 0%|          | 0.00/88.0 [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/1.12k [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/1.91M [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/65.0 [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/2.28G [00:00<?, ?B/s]
```

```
In [6]: def generate_batch_sized_chunks(list_of_elements, batch_size):
        """split the dataset into smaller batches that we can process simultaneously
        Yield successive batch-sized chunks from list_of_elements."""
        for i in range(0, len(list_of_elements), batch_size):
            yield list_of_elements[i : i + batch_size]
```

```
In [7]: def calculate_metric_on_test_ds(dataset, metric, model, tokenizer,
        batch_size=16, device=device,
        column_text="article",
```

```

column_summary="highlights"):
article_batches = list(generate_batch_sized_chunks(dataset[column_text], batch_size))
target_batches = list(generate_batch_sized_chunks(dataset[column_summary], batch_size))

for article_batch, target_batch in tqdm(
    zip(article_batches, target_batches), total=len(article_batches)):

    inputs = tokenizer(article_batch, max_length=1024, truncation=True,
        padding="max_length", return_tensors="pt")

    summaries = model.generate(input_ids=inputs["input_ids"].to(device),
        attention_mask=inputs["attention_mask"].to(device),
        length_penalty=0.8, num_beams=8, max_length=128)
    ''' parameter for length penalty ensures that the model does not generate sequences that are too long.

    # Finally, we decode the generated texts,
    # replace the token, and add the decoded texts with the references to the metric.
    decoded_summaries = [tokenizer.decode(s, skip_special_tokens=True,
        clean_up_tokenization_spaces=True)
        for s in summaries]

    decoded_summaries = [d.replace("'", " ") for d in decoded_summaries]

    metric.add_batch(predictions=decoded_summaries, references=target_batch)

# Finally compute and return the ROUGE scores.
score = metric.compute()
return score

```

Load data

Link: <https://huggingface.co/datasets/samsum>

```

In [8]: dataset_samsum = load_dataset("samsum")

split_lengths = [len(dataset_samsum[split]) for split in dataset_samsum]

print(f"Split lengths: {split_lengths}")
print(f"Features: {dataset_samsum['train'].column_names}")
print("\nDialogue:")

print(dataset_samsum["test"][1]["dialogue"])

print("\nSummary:")

print(dataset_samsum["test"][1]["summary"])

```

```

Downloading builder script: 0%|          | 0.00/3.36k [00:00<?, ?B/s]
Downloading metadata: 0%|          | 0.00/1.58k [00:00<?, ?B/s]
Downloading readme: 0%|          | 0.00/6.88k [00:00<?, ?B/s]

```

Downloading and preparing dataset samsum/samsum to /root/.cache/huggingface/datasets/samsum/samsum/0.0.0/f1d7c6b7353e6de335d444e424dc002ef70d1277109031327bc5

```

Downloading data: 0%|          | 0.00/2.94M [00:00<?, ?B/s]
Generating train split: 0%|          | 0/14732 [00:00<?, ? examples/s]
Generating test split: 0%|          | 0/819 [00:00<?, ? examples/s]
Generating validation split: 0%|          | 0/818 [00:00<?, ? examples/s]

```

Dataset samsum downloaded and prepared to /root/.cache/huggingface/datasets/samsum/samsum/0.0.0/f1d7c6b7353e6de335d444e424dc002ef70d1277109031327bc5

```

0%|          | 0/3 [00:00<?, ?it/s]

```

```

Split lengths: [14732, 819, 818]
Features: ['id', 'dialogue', 'summary']

```

```

Dialogue:
Eric: MACHINE!
Rob: That's so gr8!
Eric: I know! And shows how Americans see Russian ;)
Rob: And it's really funny!
Eric: I know! I especially like the train part!
Rob: Hahaha! No one talks to the machine like that!
Eric: Is this his only stand-up?
Rob: Idk. I'll check.
Eric: Sure.
Rob: Turns out no! There are some of his stand-ups on youtube.
Eric: Gr8! I'll watch them now!
Rob: Me too!
Eric: MACHINE!
Rob: MACHINE!
Eric: TTYL?
Rob: Sure :)

```

```

Summary:
Eric and Rob are going to watch a stand-up on youtube.

```

Evaluating PEGASUS on SAMSum

```
In [9]: dataset_samsum['test'][0]['dialogue']
```

```
Out [9]: "Hannah: Hey, do you have Betty's number?\nAmanda: Lemme check\nHannah: <file_gif>\nAmanda: Sorry, can't find it.\nAmanda: Ask  
Larry\nAmanda: He called her last time we were at the park together\nHannah: I don't know him well\nHannah: <file_gif>\nAmanda: Don't be  
shy, he's very nice\nHannah: If you say so.\nHannah: I'd rather you texted him\nAmanda: Just text him 😊\nHannah: Urgh..  
Alright\nHannah: Bye\nAmanda: Bye bye"
```

```
In [10]: pipe = pipeline('summarization', model = model_ckpt )

pipe_out = pipe(dataset_samsum['test'][0]['dialogue'] )

print(pipe_out)
```

Your max_length is set to 128, but you input_length is only 122. You might consider decreasing max_length manually, e.g. summarizer('...', max_length=61)

```
{'summary_text': "Amanda: Ask Larry Amanda: He called her last time we were at the park together .<n>Hannah: I'd rather you texted him .<n>Amanda:"
```

```
In [11]: print(pipe_out[0]['summary_text'].replace(" .<n>", ".\n"))
```

Amanda: Ask Larry Amanda: He called her last time we were at the park together.
Hannah: I'd rather you texted him.
Amanda: Just text him .

```
In [12]: rouge_metric = load_metric('rouge')

score = calculate_metric_on_test_ds(dataset_samsum['test'], rouge_metric, model_pegasus, tokenizer, column_text=
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4: FutureWarning: load_metric is deprecated and will be removed in the next major version of datasets. Use 'evaluate.load' instead, from the new library 🤗 Evaluate: <https://huggingface.co/docs/evaluate> after removing the cwd from sys.path.

Downloading builder script: 0%| | 0.00/2.16k [00:00<?, ?B/s]

```
0%|          | 0/103 [00:00<?, ?it/s]  [A
1%|          | 1/103 [00:14<24:31, 14.42s/it]  [A
2%|          | 2/103 [00:24<19:39, 11.68s/it]  [A
3%|          | 3/103 [00:36<20:10, 12.10s/it]  [A
4%|          | 4/103 [00:49<20:34, 12.47s/it]  [A
5%|          | 5/103 [00:58<18:24, 11.27s/it]  [A
6%|          | 6/103 [01:12<19:15, 11.91s/it]  [A
7%|          | 7/103 [01:25<19:40, 12.29s/it]  [A
8%|          | 8/103 [01:37<19:28, 12.30s/it]  [A
9%|          | 9/103 [01:46<17:30, 11.18s/it]  [A
10%|         | 10/103 [01:59<18:23, 11.86s/it]  [A
11%|         | 11/103 [02:09<17:09, 11.19s/it]  [A
12%|         | 12/103 [02:22<17:55, 11.82s/it]  [A
13%|         | 13/103 [02:34<17:46, 11.84s/it]  [A
14%|         | 14/103 [02:47<17:55, 12.08s/it]  [A
15%|         | 15/103 [03:00<18:14, 12.44s/it]  [A
16%|         | 16/103 [03:13<18:26, 12.72s/it]  [A
17%|         | 17/103 [03:26<18:14, 12.72s/it]  [A
17%|         | 18/103 [03:35<16:38, 11.75s/it]  [A
18%|         | 19/103 [03:46<16:00, 11.44s/it]  [A
19%|         | 20/103 [03:59<16:35, 12.00s/it]  [A
20%|         | 21/103 [04:13<16:54, 12.37s/it]  [A
21%|         | 22/103 [04:24<16:11, 11.99s/it]  [A
22%|         | 23/103 [04:33<15:01, 11.27s/it]  [A
23%|         | 24/103 [04:47<15:38, 11.88s/it]  [A
24%|         | 25/103 [04:58<15:21, 11.81s/it]  [A
25%|         | 26/103 [05:11<15:33, 12.12s/it]  [A
26%|         | 27/103 [05:25<15:49, 12.49s/it]  [A
27%|         | 28/103 [05:36<15:02, 12.04s/it]  [A
28%|         | 29/103 [05:45<13:57, 11.32s/it]  [A
29%|         | 30/103 [05:59<14:31, 11.94s/it]  [A
30%|         | 31/103 [06:11<14:25, 12.03s/it]  [A
31%|         | 32/103 [06:22<13:59, 11.82s/it]  [A
32%|         | 33/103 [06:35<14:19, 12.27s/it]  [A
33%|         | 34/103 [06:46<13:24, 11.66s/it]  [A
34%|         | 35/103 [06:54<11:59, 10.57s/it]  [A
35%|         | 36/103 [07:06<12:25, 11.13s/it]  [A
36%|         | 37/103 [07:19<12:57, 11.78s/it]  [A
37%|         | 38/103 [07:33<13:17, 12.27s/it]  [A
38%|         | 39/103 [07:46<13:25, 12.59s/it]  [A
39%|         | 40/103 [07:59<13:12, 12.57s/it]  [A
40%|         | 41/103 [08:08<12:06, 11.71s/it]  [A
41%|         | 42/103 [08:19<11:34, 11.38s/it]  [A
42%|         | 43/103 [08:32<11:58, 11.97s/it]  [A
43%|         | 44/103 [08:40<10:22, 10.55s/it]  [A
44%|         | 45/103 [08:53<10:59, 11.37s/it]  [A
45%|         | 46/103 [09:02<10:03, 10.59s/it]  [A
46%|         | 47/103 [09:14<10:29, 11.24s/it]  [A
47%|         | 48/103 [09:23<09:33, 10.42s/it]  [A
48%|         | 49/103 [09:32<09:06, 10.13s/it]  [A
49%|         | 50/103 [09:46<09:48, 11.10s/it]  [A
50%|         | 51/103 [09:56<09:16, 10.71s/it]  [A
50%|         | 52/103 [10:09<09:45, 11.47s/it]  [A
51%|         | 53/103 [10:21<09:52, 11.84s/it]  [A
52%|         | 54/103 [10:33<09:37, 11.78s/it]  [A
53%|         | 55/103 [10:46<09:47, 12.24s/it]  [A
54%|         | 56/103 [10:58<09:25, 12.02s/it]  [A
55%|         | 57/103 [11:10<09:10, 11.97s/it]  [A
56%|         | 58/103 [11:20<08:39, 11.55s/it]  [A
57%|         | 59/103 [11:34<08:50, 12.06s/it]  [A
58%|         | 60/103 [11:47<08:54, 12.42s/it]  [A
59%|         | 61/103 [11:57<08:12, 11.73s/it]  [A
60%|         | 62/103 [12:09<08:02, 11.77s/it]  [A
61%|         | 63/103 [12:21<07:52, 11.81s/it]  [A
```

62%		64/103	[12:34<07:57, 12.25s/it]	0[A
63%		65/103	[12:47<07:55, 12.51s/it]	0[A
64%		66/103	[12:58<07:21, 11.93s/it]	0[A
65%		67/103	[13:08<06:53, 11.48s/it]	0[A
66%		68/103	[13:21<06:56, 11.91s/it]	0[A
67%		69/103	[13:30<06:17, 11.10s/it]	0[A
68%		70/103	[13:43<06:21, 11.57s/it]	0[A
69%		71/103	[13:55<06:12, 11.65s/it]	0[A
70%		72/103	[14:08<06:16, 12.15s/it]	0[A
71%		73/103	[14:17<05:38, 11.29s/it]	0[A
72%		74/103	[14:30<05:35, 11.58s/it]	0[A
73%		75/103	[14:40<05:12, 11.18s/it]	0[A
74%		76/103	[14:53<05:19, 11.83s/it]	0[A
75%		77/103	[15:06<05:18, 12.26s/it]	0[A
76%		78/103	[15:20<05:14, 12.57s/it]	0[A
77%		79/103	[15:33<05:07, 12.80s/it]	0[A
78%		80/103	[15:46<04:57, 12.95s/it]	0[A
79%		81/103	[15:59<04:45, 12.97s/it]	0[A
80%		82/103	[16:13<04:34, 13.06s/it]	0[A
81%		83/103	[16:22<03:56, 11.82s/it]	0[A
82%		84/103	[16:35<03:53, 12.27s/it]	0[A
83%		85/103	[16:48<03:46, 12.58s/it]	0[A
83%		86/103	[16:59<03:26, 12.14s/it]	0[A
84%		87/103	[17:13<03:19, 12.49s/it]	0[A
85%		88/103	[17:26<03:09, 12.61s/it]	0[A
86%		89/103	[17:39<02:59, 12.80s/it]	0[A
87%		90/103	[17:51<02:44, 12.68s/it]	0[A
88%		91/103	[18:03<02:28, 12.38s/it]	0[A
89%		92/103	[18:16<02:19, 12.66s/it]	0[A
90%		93/103	[18:30<02:08, 12.85s/it]	0[A
91%		94/103	[18:42<01:53, 12.62s/it]	0[A
92%		95/103	[18:50<01:31, 11.45s/it]	0[A
93%		96/103	[19:04<01:23, 12.00s/it]	0[A
94%		97/103	[19:17<01:14, 12.38s/it]	0[A
95%		98/103	[19:30<01:03, 12.66s/it]	0[A
96%		99/103	[19:43<00:51, 12.85s/it]	0[A
97%		100/103	[19:56<00:38, 12.81s/it]	0[A
98%		101/103	[20:09<00:25, 12.74s/it]	0[A
99%		102/103	[20:22<00:12, 12.91s/it]	0[A
100%		103/103	[20:25<00:00, 11.90s/it]	

```
In [13]: rouge_names = ["rouge1", "rouge2", "rougeL", "rougeLsum"]
rouge_dict = dict((rn, score[rn].mid.fmeasure ) for rn in rouge_names )

pd.DataFrame(rouge_dict, index = ['pegasus'])
```

```
Out [13]:
```

	rouge1	rouge2	rougeL	rougeLsum
pegasus	0.015558	0.000295	0.015537	0.015554


Histogram

```
In [ ]: dialogue_token_len = len([tokenizer.encode(s) for s in dataset_samsum['train']['dialogue']])

summary_token_len = len([tokenizer.encode(s) for s in dataset_samsum['train']['summary']])

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
axes[0].hist(dialogue_token_len, bins = 20, color = 'C0', edgecolor = 'C0' )
axes[0].set_title("Dialogue Token Length")
axes[0].set_xlabel("Length")
axes[0].set_ylabel("Count")

axes[1].hist(summary_token_len, bins = 20, color = 'C0', edgecolor = 'C0' )
axes[1].set_title("Summary Token Length")
axes[1].set_xlabel("Length")
plt.tight_layout()
plt.show()
```



```
In [14]: def convert_examples_to_features(example_batch):
input_encodings = tokenizer(example_batch['dialogue'] , max_length = 1024, truncation = True )
```

```

with tokenizer.as_target_tokenizer():
    target_encodings = tokenizer(example_batch['summary'], max_length = 128, truncation = True )

    return {
        'input_ids' : input_encodings['input_ids'],
        'attention_mask': input_encodings['attention_mask'],
        'labels': target_encodings['input_ids']
    }

dataset_samsum_pt = dataset_samsum.map(convert_examples_to_features, batched = True)

```

```
0%|          | 0/15 [00:00<?, ?ba/s]
```

```

/usr/local/lib/python3.7/dist-packages/transformers/tokenization_utils_base.py:3547: UserWarning: `as_target_tokenizer` is deprecated and will be removed in v5 of Transformers. You can tokenize your labels by using the argument `text_target` of the regular `__call__` method (either in the same call as your input texts if you use the same keyword arguments, or in a separate call.
  "as_target_tokenizer` is deprecated and will be removed in v5 of Transformers. You can tokenize your "

```

```

0%|          | 0/1 [00:00<?, ?ba/s]
0%|          | 0/1 [00:00<?, ?ba/s]

```

In [15]: `from transformers import DataCollatorForSeq2Seq`

```
seq2seq_data_collator = DataCollatorForSeq2Seq(tokenizer, model=model_pegasus)
```

In [17]: `from google.colab import drive`
`drive.mount('/content/drive')`

Mounted at /content/drive

In [18]: `%cd /content/drive/MyDrive/005_BOKTIAR_AHMED_BAPPY/My_classes/FSDS_NOV_10_AM`

/content/drive/MyDrive/005_BOKTIAR_AHMED_BAPPY/My_classes/FSDS_NOV_10_AM

In [19]: `from transformers import TrainingArguments, Trainer`

```

trainer_args = TrainingArguments(
    output_dir='pegasus-samsum', num_train_epochs=1, warmup_steps=500,
    per_device_train_batch_size=1, per_device_eval_batch_size=1,
    weight_decay=0.01, logging_steps=10,
    evaluation_strategy='steps', eval_steps=500, save_steps=1e6,
    gradient_accumulation_steps=16
)

```

In [20]: `trainer = Trainer(model=model_pegasus, args=trainer_args,`
`tokenizer=tokenizer, data_collator=seq2seq_data_collator,`
`train_dataset=dataset_samsum_pt["train"],`
`eval_dataset=dataset_samsum_pt["validation"])`

In [21]: `trainer.train()`

```

The following columns in the training set don't have a corresponding argument in `PegasusForConditionalGeneration.forward` and have been ignored: summary, dialogue, id. If summary, dialogue, id are not expected by `PegasusForConditionalGeneration.forward`, you can safely ignore this message.
/usr/local/lib/python3.7/dist-packages/transformers/optimization.py:310: FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Use the PyTorch implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True` to disable this warning
  FutureWarning,
**** Running training ****
Num examples = 14732
Num Epochs = 1
Instantaneous batch size per device = 1
Total train batch size (w. parallel, distributed & accumulation) = 16
Gradient Accumulation steps = 16
Total optimization steps = 920
Number of trainable parameters = 568699904
You're using a PegasusTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.

```

[920/920 45:04, Epoch 0/1]

Step	Training Loss	Validation Loss
------	---------------	-----------------

500	1.690200	1.488429
-----	----------	----------

```

The following columns in the evaluation set don't have a corresponding argument in `PegasusForConditionalGeneration.forward` and have been ignored: summary, dialogue, id. If summary, dialogue, id are not expected by `PegasusForConditionalGeneration.forward`, you can safely ignore this message.
**** Running Evaluation ****
Num examples = 818
Batch size = 1

```

Training completed. Do not forget to share your model on huggingface.co/models =)

Out [21]: TrainOutput(global_step=920, training_loss=1.828826087454091, metrics={'train_runtime': 2707.8434, 'train_samples_per_second': 5.44, 'train_steps_per_second': 0.34, 'total_flos': 5526961323663360.0, 'train_loss': 1.828826087454091, 'epoch': 1.0})

```
In [30]: score = calculate_metric_on_test_ds(
        dataset_samsum['test'], rouge_metric, trainer.model, tokenizer, batch_size = 2, column_text = 'dialogue',
    )

rouge_dict = dict((rn, score[rn].mid.fmeasure ) for rn in rouge_names )

pd.DataFrame(rouge_dict, index = [f'pegasus'] )
```

100%|██████████| 410/410 [13:08<00:00, 1.92s/it]

```
Out [30]:
```

	rouge1	rouge2	rougeL	rougeLsum
pegasus	0.018618	0.000297	0.018493	0.018518

```
In [22]: ## Save model
model_pegasus.save_pretrained("pegasus-samsum-model")
```

Configuration saved in pegasus-samsum-model/config.json
Model weights saved in pegasus-samsum-model/pytorch_model.bin

```
In [23]: ## Save tokenizer
tokenizer.save_pretrained("tokenizer")
```

tokenizer config file saved in tokenizer/tokenizer_config.json
Special tokens file saved in tokenizer/special_tokens_map.json

```
Out [23]: ('tokenizer/tokenizer_config.json',
'tokenizer/special_tokens_map.json',
'tokenizer/spiece.model',
'tokenizer/added_tokens.json',
'tokenizer/tokenizer.json')
```

Test

```
In [24]: dataset_samsum = load_dataset("samsum")
```

WARNING:datasets.builder:Found cached dataset samsum
(/root/.cache/huggingface/datasets/samsum/samsum/0.0.0/f1d7c6b7353e6de335d444e424dc002ef70d1277109031327bc9cc6af5d3d46e)

0%| | 0/3 [00:00<?, ?it/s]

```
In [25]: tokenizer = AutoTokenizer.from_pretrained("tokenizer")
```

loading file spiece.model
loading file tokenizer.json
loading file added_tokens.json
loading file special_tokens_map.json
loading file tokenizer_config.json

```
In [27]: sample_text = dataset_samsum["test"][0]["dialogue"]

reference = dataset_samsum["test"][0]["summary"]
```

```
In [28]: gen_kwargs = {"length_penalty": 0.8, "num_beams":8, "max_length": 128}

pipe = pipeline("summarization", model="pegasus-samsum-model",tokenizer=tokenizer)
```

loading configuration file pegasus-samsum-model/config.json

```
Model config PegasusConfig {
  "name_or_path": "pegasus-samsum-model",
  "activation_dropout": 0.1,
  "activation_function": "relu",
  "add_bias_logits": false,
  "add_final_layer_norm": true,
  "architectures": [
    "PegasusForConditionalGeneration"
  ],
  "attention_dropout": 0.1,
  "bos_token_id": 0,
  "classif_dropout": 0.0,
  "classifier_dropout": 0.0,
  "d_model": 1024,
  "decoder_attention_heads": 16,
  "decoder_ffn_dim": 4096,
  "decoder_layerdrop": 0.0,
  "decoder_layers": 16,
  "decoder_start_token_id": 0,
  "dropout": 0.1,
  "encoder_attention_heads": 16,
  "encoder_ffn_dim": 4096,
  "encoder_layerdrop": 0.0,
  "encoder_layers": 16,
  "eos_token_id": 1,
  "extra_pos_embeddings": 1,
  "forced_eos_token_id": 1,
  "id2label": {
    "0": "LABEL_0",
    "1": "LABEL_1",
    "2": "LABEL_2"
  },
  "init_std": 0.02,
  "is_encoder_decoder": true,
```

```

"label2id": {
  "LABEL_0": 0,
  "LABEL_1": 1,
  "LABEL_2": 2
},
"length_penalty": 0.8,
"max_length": 128,
"max_position_embeddings": 1024,
"min_length": 32,
"model_type": "pegasus",
"normalize_before": true,
"normalize_embedding": false,
"num_beams": 8,
"num_hidden_layers": 16,
"pad_token_id": 0,
"scale_embedding": true,
"static_position_embeddings": true,
"torch_dtype": "float32",
"transformers_version": "4.24.0",
"use_cache": true,
"vocab_size": 96103
}

```

loading configuration file pegasus-samsum-model/config.json

```

Model config PegasusConfig {
  "_name_or_path": "pegasus-samsum-model",
  "activation_dropout": 0.1,
  "activation_function": "relu",
  "add_bias_logits": false,
  "add_final_layer_norm": true,
  "architectures": [
    "PegasusForConditionalGeneration"
  ],
  "attention_dropout": 0.1,
  "bos_token_id": 0,
  "classif_dropout": 0.0,
  "classifier_dropout": 0.0,
  "d_model": 1024,
  "decoder_attention_heads": 16,
  "decoder_ffn_dim": 4096,
  "decoder_layerdrop": 0.0,
  "decoder_layers": 16,
  "decoder_start_token_id": 0,
  "dropout": 0.1,
  "encoder_attention_heads": 16,
  "encoder_ffn_dim": 4096,
  "encoder_layerdrop": 0.0,
  "encoder_layers": 16,
  "eos_token_id": 1,
  "extra_pos_embeddings": 1,
  "forced_eos_token_id": 1,
  "id2label": {
    "0": "LABEL_0",
    "1": "LABEL_1",
    "2": "LABEL_2"
  },
  "init_std": 0.02,
  "is_encoder_decoder": true,
  "label2id": {
    "LABEL_0": 0,
    "LABEL_1": 1,
    "LABEL_2": 2
  },
  "length_penalty": 0.8,
  "max_length": 128,
  "max_position_embeddings": 1024,
  "min_length": 32,
  "model_type": "pegasus",
  "normalize_before": true,
  "normalize_embedding": false,
  "num_beams": 8,
  "num_hidden_layers": 16,
  "pad_token_id": 0,
  "scale_embedding": true,
  "static_position_embeddings": true,
  "torch_dtype": "float32",
  "transformers_version": "4.24.0",
  "use_cache": true,
  "vocab_size": 96103
}

```

loading weights file pegasus-samsum-model/pytorch_model.bin

All model checkpoint weights were used when initializing PegasusForConditionalGeneration.

All the weights of PegasusForConditionalGeneration were initialized from the model checkpoint at pegasus-samsum-model.

If your task is similar to the task the model of the checkpoint was trained on, you can already use PegasusForConditionalGeneration for predictions without further training.

```

In [29]: print("Dialogue:")
         print(sample_text)

         print("\nReference Summary:")
         print(reference)

         print("\nModel Summary:")
         print(pipe(sample_text, **gen_kwargs)[0]["summary_text"])

```

Your max_length is set to 128, but you input_length is only 122. You might consider decreasing max_length manually, e.g. summarizer('...', max_length=61)

```

Dialogue:
Hannah: Hey, do you have Betty's number?
Amanda: Lemme check
Hannah: <file_gif>
Amanda: Sorry, can't find it.
Amanda: Ask Larry
Amanda: He called her last time we were at the park together
Hannah: I don't know him well
Hannah: <file_gif>
Amanda: Don't be shy, he's very nice
Hannah: If you say so..

```

Hannah: I'd rather you texted him
Amanda: Just text him 😊
Hannah: Urgh.. Alright
Hannah: Bye
Amanda: Bye bye

Reference Summary:

Hannah needs Betty's number but Amanda doesn't have it. She needs to contact Larry.

Model Summary:

Amanda can't find Betty's number. Larry called Betty last time they were at the park together. Hannah wants Amanda to text Larry. Amanda will text L